

Distributed Pharmacy Inventory Management System

Shaoming Pan
Chenchuhui Hu

gRPC Microservice Architecture vs REST Monolith

Agenda

01

System Requirements

6 functional + non-functional requirements

04

Evaluation Methodology

Benchmark setup and workload specification

02

Architecture Design

gRPC Microservice (6 nodes) + REST Monolith

05

Performance Results

Latency & throughput comparison

03

Communication Model

gRPC with Protobuf vs HTTP/JSON

06

AI Tools & Lessons

How Claude accelerated development

System Requirements



Add Drug

Register new drug with name, quantity, price, expiry



Get Drug

Query drug details by ID



Update Stock

Modify inventory quantity



Delete Drug

Remove drug record from system



List All Drugs


View complete inventory list



Low Stock Alert

Find drugs below threshold quantity

Architecture Design — gRPC Microservice (6 Nodes)



Communication Model

gRPC Microservice

Protocol: HTTP/2 + Protobuf binary

Format: Binary serialization

Speed: Faster encoding/decoding

Streaming: Bidirectional streams

Contract: Strongly typed .proto file

Best for: Internal service-to-service calls

REST Monolith

Protocol: HTTP/1.1 + JSON text

Format: Human-readable JSON

Speed: Slower JSON parsing

Streaming: Not supported

Contract: Flexible, no schema required

Best for: Public APIs, browser clients

Evaluation Methodology



Hardware Environment

Machine: MacBook Pro (Apple Silicon)

RAM: 16 GB | OS: macOS

Docker Desktop (all containers local)

Python 3.9 benchmark scripts



Workload Specification

Write: concurrent AddDrug requests

Read: concurrent ListDrugs requests

Concurrency: 10, 50, 100, 500, 1000

Metrics: Avg Latency (ms) + Throughput (req/s)

$$\text{Latency} = T_{\text{response}} - T_{\text{request}}$$


Time

$$\text{Throughput} = \frac{\text{Total Requests}}{\text{Total Time}}$$


gRPC: 6 nodes | REST Monolith: 2 containers | Both running simultaneously

Performance Results

Write Throughput (req/s)



Read Latency (ms) — Lower is Better



At 1000 users: gRPC read latency 1023ms vs REST 5150ms — 5x better | gRPC write throughput 63% higher

Trade-off Analysis

Dimension	gRPC Microservice	REST Monolith
Performance	Higher throughput, lower latency	Degrades sharply at scale
Scalability	Horizontal via load balancer	Single bottleneck
Fault Tolerance	Redundant servers + DB replication	Single point of failure
Complexity	Higher — multi-node orchestration	Lower — one container
Dev Speed	Slower — proto definitions needed	Faster — standard REST

AI Tools — How Claude Helped



Architecture Design

Designed the 6-node system layout and Docker Compose configuration



Proto File

Generated all gRPC .proto definitions for 6 service methods



Debugging

Diagnosed Dockerfile COPY path error and fixed build context



Benchmarking

Generated concurrent Python benchmark and plotting scripts



Documentation

Wrote README, final report, and this presentation

Prompting Strategy: provide exact errors + architecture constraints → targeted, working solutions

Conclusion

-  6 functional requirements fully implemented and tested
-  6-node gRPC microservice architecture deployed with Docker
-  gRPC delivers 5× lower latency and 63% higher throughput at scale
-  Microservice design provides fault tolerance REST cannot match
-  AI tools (Claude) accelerated development at every stage